Person 1: Measured Results

Interlock V2 | 30 August 2026

Executive result

The offline F-019 evaluation compared 33 decision strategies on three independently seeded 200-case sets and all 300 labelled anchors. The selected calibrated risk floor records no false interventions on the measured clean cases while preserving all induced-defect catches. Net spend remains below its release target, and the finite clean sample is not large enough to prove a production rate below 2% with 95% confidence.

Measure	Result
Policy	banking-v3@sha256:4b63a6b7b5416683
Production strategy	1% risk floor + 50% relative gain
Seeded evaluations	3 x 200 fixed-answer cases
Anchor records	300 records / 300 unique IDs
Calibration set	10,000 document-disjoint examples

System path

```
User -> Gateway -> stakes/routing -> Qwen tier
                         -> commit gate -> observer/risk engine
                         -> pass | repair | reroute | hold | block
                         -> ConsoleHub events + ledger economics
```

300-anchor exhaustive checks

Check	Result
Record count	300
Unique item IDs	300
Failure-mode distribution	clause_swapped: 6, clean: 270, contradiction: 5, number_corrupted: 7, retrieval_dropped: 7, unanswerable: 5
Clean answers contained in supplied context	270/270
Label consistency errors	0

Interpretation: these checks establish dataset completeness and internal consistency. The anchor policy result is 0/270 false interventions and 30/30 defect catches. Because the anchors participate in calibration, they are supporting evidence rather than an independent holdout, and they do not prove post-action efficacy.

Measured system evidence

Area	Three-seed result	Status
Pre-action catch rate	100.00% / 100.00% / 100.00%	PASS
Added p95 latency	16 ms worst seed	PASS
Verification cost	0.78% / 0.48% / 0.62%	PASS
Net spend change	-22.56% / -20.67% / -20.81%	MISS: target about -30%
Ungrounded escapes	0.00% / 0.00% / 0.00%	PASS empirically
False interventions	0/140 on each seed	PASS point estimate
False-intervention 95% CI	0-2.67%	Needs larger clean holdout

Calibration and decision policy

Measurement	Result
Out-of-fold ECE	0.0046
Brier score	0.0005
Induced-taxonomy AUROC	1.0000
Mean clean calibrated risk	0.2568%
Production risk floor	1.00%
Conformal threshold / intervention	37.00% / 9.22%

Why this method was selected

The old question-drift signal compared question words directly with terse answer words and was effectively inverted. It now compares both through retrieved context. Numeric corruption generation also retries when a replacement figure remains supported elsewhere in the passage. After recalibration, clean risk averages 0.2568% while each induced defect mode averages at least 94.88%. A 1% floor therefore separates the measured clean and defective cases without changing the rupee impact model. Deterministic hard rules, conformal filtering, and unavailable-action constraints still take precedence.

Claim boundary

The evidence supports a measured false-intervention result of 0/140 per seed, not a claim of zero production false interventions. The Wilson upper bound is 2.67%, above the 2% target. The induced taxonomy is highly separable and must be challenged with naturally occurring defects and a larger untouched clean set. Reviewed forced-action outcomes are also still required before claiming that repair, hold, or block improves final answers.

Conclusion

Within the offline paired evaluation, F-019 is resolved: the selected policy removes all measured false interventions without a measured catch or escape regression. Verification cost also passes. Net spend remains the only scorecard miss, and statistical confidence plus external-defect validity remain explicit release evidence requirements.