

Person 1: Measured Results

Interlock V2 | 29 August 2026

Executive result

The core Person 1 implementation is present and passes the repository verification suite. The 300-item anchor set has been exhaustively checked for integrity and policy-label consistency. Live Ollama is now available with Qwen 4B and Qwen 8B, but CPU inference is too slow for an honest full multi-model judge run in this session. The report therefore separates measured results from open evidence gates.

Measure	Result
Repository tests	906 passed
Chaos tests	7 passed
Ruff	Clean
Anchor records	300 records / 300 unique IDs
Policy hash	banking-v3@sha256:0aeb8a4b17fe218c

System path

```
User -> Gateway -> stakes/routing -> Qwen tier
                                -> commit gate -> observer/risk engine
                                -> pass | repair | reroute | hold | block
                                -> ConsoleHub events + ledger economics
```

300-anchor exhaustive checks

Check	Result
Record count	300
Unique item IDs	300
Failure-mode distribution	clause_swapped: 6, clean: 270, contradiction: 5, number_corrupted: 7, retrieval_dropped: 7, unanswerable: 5
Clean answers contained in supplied context	270/270
Label consistency errors	0

Interpretation: these checks prove that the reference dataset is complete, unique, document-backed, and internally consistent. They do not prove that an intervention fixed an answer.

Measured system evidence

Area	Measured result	Status
Pre-action catch rate	100.00% [91.8, 100.0]	PASS
Added p95 latency	15 ms	PASS
Verification cost	5.20%	MISS: target <=5%
Net spend change	-18.96%	MISS: target about -30%
Ungrounded escapes	0.00% empirical	PASS; certified bound separate
False interventions	85.35%	MISS: target <=2%
Load pass	4,023 requests / 20 concurrency / 0 failures	PASS; gateway p95 531 ms
Observer warm p95	37.06 ms	MISS: target 25 ms

Live model inventory

Model	Format	Role
qwen3:4b	Q4_K_M, 4.0B	cheap/low-stakes tier
qwen3:8b	Q4_K_M, 8.2B	strong/high-stakes tier

Automated judge run

A reproducible judge runner was added at `scripts/eval_manual_anchors.py`. It records the model, raw judgment, parsed label, confidence/rationale, agreement, and latency. A short Qwen 4B run completed 4 rows before CPU throughput became impractical (about 33 seconds per row). It is retained as a smoke sample only and is not presented as a 300-item result.

Open gates

1. Run a complete 300-item judge pass when adequate compute is available. 2. Generate reviewed post-action outcomes for L1-L5; pre-action labels cannot establish intervention efficacy. 3. Resolve F-019, because the current impact model produces the measured false-intervention miss. 4. Complete live-model rehearsal, true KV-cache performance validation, and clean-checkout deployment proof.

Conclusion

Person 1's core code and operational wiring are verified. The evidence supports claiming detection, routing, event publication, holds, economics surfaces, and deterministic end-to-end operation. It does not support claiming that every intervention improves answers or that all six target metrics pass. Those claims require the open measurements above.